

Waze Project

Course 2 - Get Started with Python

Welcome to the Waze Project!

Your Waze data analytics team is still in the early stages of their user churn project. Previously, you were asked to complete a project proposal by your supervisor, May Santner. You have received notice that your project proposal has been approved and that your team has been given access to Waze's user data. To get clear insights, the user data must be inspected and prepared for the upcoming process of exploratory data analysis (EDA).

A Python notebook has been prepared to guide you through this project. Answer the questions and create an executive summary for the Waze data team.

Course 2 End-of-course project: Inspect and analyze data

In this activity, you will examine data provided and prepare it for analysis. This activity will help ensure the information is,

1. Ready to answer questions and yield insights
2. Ready for visualizations
3. Ready for future hypothesis testing and statistical methods

The purpose of this project is to investigate and understand the data provided.

The goal is to use a dataframe constructed within Python, perform a cursory inspection of the provided dataset, and inform team members of your findings.

This activity has three parts:

Part 1: Understand the situation

- How can you best prepare to understand and organize the provided information?

Part 2: Understand the data

- Create a pandas dataframe for data learning, future exploratory data analysis (EDA), and statistical activities
- Compile summary information about the data to inform next steps

Part 3: Understand the variables

- Use insights from your examination of the summary data to guide deeper investigation into variables

Follow the instructions and answer the following questions to complete the activity. Then, you will complete an Executive Summary using the questions listed on the PACE Strategy Document.

Be sure to complete this activity before moving on. The next course item will provide you with a completed exemplar to compare to your own work.

Identify data types and compile summary information

Task 1. Understand the situation

- How can you best prepare to understand and organize the provided driver data?

Begin by exploring your dataset and consider reviewing the Data Dictionary.

Task 2a. Imports and data loading

Start by importing the packages that you will need to load and explore the dataset. Make sure to use the following import statements:

- `import pandas as pd`
- `import numpy as np`

```
In [ ]: # Import packages for data manipulation
import pandas as pd
import numpy as np
```

Then, load the dataset into a dataframe. Creating a dataframe will help you conduct data manipulation, exploratory data analysis (EDA), and statistical activities.

Note: As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset

and proceed with this lab. Please continue with this activity by completing the following instructions.

```
In [2]: # Load dataset into dataframe
df = pd.read_csv('waze_dataset.csv')
```

Task 2b. Summary information

View and inspect summary information about the dataframe by **coding the following**:

1. `df.head(10)`
2. `df.info()`

Consider the following questions:

1. When reviewing the `df.head()` output, are there any variables that have missing values?
2. When reviewing the `df.info()` output, what are the data types? How many rows and columns do you have?
3. Does the dataset have any missing values?

```
In [3]: ### YOUR CODE HERE ###
df.head(10)
```

```
Out[3]:
```

	ID	label	sessions	drives	total_sessions	n_days_after_onboarding	total_navig
0	0	retained	283	226	296.748273		2276
1	1	retained	133	107	326.896596		1225
2	2	retained	114	95	135.522926		2651
3	3	retained	49	40	67.589221		15
4	4	retained	84	68	168.247020		1562
5	5	retained	113	103	279.544437		2637
6	6	retained	3	2	236.725314		360
7	7	retained	39	35	176.072845		2999
8	8	retained	57	46	183.532018		424
9	9	churned	84	68	244.802115		2997

```
In [5]: ### YOUR CODE HERE ###
df.info()
df.describe()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ID                                     14999 non-null  int64
1   label                                14299 non-null  object
2   sessions                             14999 non-null  int64
3   drives                               14999 non-null  int64
4   total_sessions                       14999 non-null  float64
5   n_days_after_onboarding              14999 non-null  int64
6   total_navigations_fav1               14999 non-null  int64
7   total_navigations_fav2               14999 non-null  int64
8   driven_km_drives                     14999 non-null  float64
9   duration_minutes_drives               14999 non-null  float64
10  activity_days                         14999 non-null  int64
11  driving_days                         14999 non-null  int64
12  device                               14999 non-null  object
dtypes: float64(3), int64(8), object(2)
memory usage: 1.5+ MB

```

Out[5]:

	ID	sessions	drives	total_sessions	n_days_after_onbo
count	14999.000000	14999.000000	14999.000000	14999.000000	14999.0
mean	7499.000000	80.633776	67.281152	189.964447	1749.8
std	4329.982679	80.699065	65.913872	136.405128	1008.5
min	0.000000	0.000000	0.000000	0.220211	4.0
25%	3749.500000	23.000000	20.000000	90.661156	878.0
50%	7499.000000	56.000000	48.000000	159.568115	1741.0
75%	11248.500000	112.000000	93.000000	254.192341	2623.5
max	14998.000000	743.000000	596.000000	1216.154633	3500.0

Task 2c. Null values and summary statistics

Compare the summary statistics of the 700 rows that are missing labels with summary statistics of the rows that are not missing any values.

Question: Is there a discernible difference between the two populations?

```

In [7]: # Isolate rows with null values
null_rows = df[df.isnull().any(axis=1)]
print(null_rows)

# Display summary stats of rows with null values
### YOUR CODE HERE ###
null_rows.describe()

```

	ID	label	sessions	drives	total_sessions	n_days_after_onboardin
g \						
77	77	NaN	63	50	133.104155	78
3						
80	80	NaN	116	93	436.060183	158
4						
98	98	NaN	78	64	583.492789	341
4						
111	111	NaN	106	102	113.379056	222
8						
142	142	NaN	32	26	222.129310	20
8						
...	
...						
14941	14941	NaN	191	160	485.328204	128
7						
14943	14943	NaN	48	38	96.797017	55
5						
14945	14945	NaN	34	29	134.416604	164
3						
14972	14972	NaN	220	181	256.212166	171
8						
14993	14993	NaN	67	57	97.570074	113
1						

	total_navigations_fav1	total_navigations_fav2	driven_km_drives	\
77	201	0	2649.015822	
80	283	62	4183.409514	
98	0	0	1811.140893	
111	14	0	2817.481840	
142	55	10	2459.816477	
...	
14941	25	0	6468.181924	
14943	0	6	8266.129497	
14945	268	2	4554.007843	
14972	360	23	5586.913459	
14993	207	102	2267.052913	

	duration_minutes_drives	activity_days	driving_days	device
77	1517.209970	19	13	iPhone
80	3121.889952	18	15	iPhone
98	642.189122	12	11	Android
111	2011.724274	17	13	Android
142	874.427617	11	7	iPhone
...
14941	3466.104564	14	14	iPhone
14943	5902.351711	19	19	iPhone
14945	1579.211201	18	17	Android
14972	4104.440202	19	18	iPhone
14993	318.120634	27	26	iPhone

[700 rows x 13 columns]

Out [7]:

	ID	sessions	drives	total_sessions	n_days_after_onboardin
count	700.000000	700.000000	700.000000	700.000000	700.000000
mean	7405.584286	80.837143	67.798571	198.483348	1709.2957
std	4306.900234	79.987440	65.271926	140.561715	1005.30656
min	77.000000	0.000000	0.000000	5.582648	16.000000
25%	3744.500000	23.000000	20.000000	94.056340	869.000000
50%	7443.000000	56.000000	47.500000	177.255925	1650.500000
75%	11007.000000	112.250000	94.000000	266.058022	2508.750000
max	14993.000000	556.000000	445.000000	1076.879741	3498.000000

In [9]:

```
# Isolate rows without null values
non_null_rows = df.dropna()

# Display summary stats of rows without null values
non_null_rows.describe()
```

Out [9]:

	ID	sessions	drives	total_sessions	n_days_after_onbo
count	14299.000000	14299.000000	14299.000000	14299.000000	14299.0
mean	7503.573117	80.623820	67.255822	189.547409	1751.8
std	4331.207621	80.736502	65.947295	136.189764	1008.6
min	0.000000	0.000000	0.000000	0.220211	4.0
25%	3749.500000	23.000000	20.000000	90.457733	878.5
50%	7504.000000	56.000000	48.000000	158.718571	1749.0
75%	11257.500000	111.000000	93.000000	253.540450	2627.5
max	14998.000000	743.000000	596.000000	1216.154633	3500.0

The Number of null values are nearly half of the nonnull values.

Task 2d. Null values - device counts

Next, check the two populations with respect to the `device` variable.

Question: How many iPhone users had null values and how many Android users had null values?

In [16]:

```
# Get count of null values by device
### YOUR CODE HERE ###
```

```
device_count = null_rows['device'].nunique()
```

700

Now, of the rows with null values, calculate the percentage with each device—Android and iPhone. You can do this directly with the `value_counts()` function.

```
In [19]: # Get rows with null values
null_rows = df[df.isnull().any(axis=1)]
# Calculate % of iPhone nulls and Android nulls
total_iphone_rows = df[df['device'] == 'iPhone'].shape[0]
total_android_rows = df[df['device'] == 'Android'].shape[0]
iphone_null_percentage = (null_rows[null_rows['device'] == 'iPhone'].shape[0] / total_iphone_rows) * 100
android_null_percentage = (null_rows[null_rows['device'] == 'Android'].shape[0] / total_android_rows) * 100

print("Percentage of null values for iPhone:", iphone_null_percentage)
print("Percentage of null values for Android:", android_null_percentage)
```

```
Percentage of null values for iPhone: 4.621588089330024
Percentage of null values for Android: 4.749389900506852
```

How does this compare to the device ratio in the full dataset?

```
In [20]: # Calculate % of iPhone users and Android users in full dataset
# Calculate the device ratio in the full dataset
total_devices = df.shape[0]
iphone_ratio = (df[df['device'] == 'iPhone'].shape[0] / total_devices) * 100
android_ratio = (df[df['device'] == 'Android'].shape[0] / total_devices) * 100

print("Percentage of null values for iPhone:", iphone_null_percentage)
print("Percentage of null values for Android:", android_null_percentage)
print("Device ratio in the full dataset:")
print("iPhone:", iphone_ratio)
print("Android:", android_ratio)
```

```
Percentage of null values for iPhone: 4.621588089330024
Percentage of null values for Android: 4.749389900506852
Device ratio in the full dataset:
iPhone: 64.48429895326355
Android: 35.515701046736446
```

The percentage of missing values by each device is consistent with their representation in the data overall.

There is nothing to suggest a non-random cause of the missing data.

Examine the counts and percentages of users who churned vs. those who were retained. How many of each group are represented in the data?

```
In [25]: # Calculate counts of churned vs. retained
churned_count = df[df['label'] == 'churned'].shape[0]
retained_count = df[df['label'] == 'retained'].shape[0]
```

```
print("Churned count:", churned_count)
print("Retained count:", retained_count)
```

Churned count: 2536
Retained count: 11763

This dataset contains 82% retained users and 18% churned users.

Next, compare the medians of each variable for churned and retained users. The reason for calculating the median and not the mean is that you don't want outliers to unduly affect the portrayal of a typical user. Notice, for example, that the maximum value in the `driven_km_drives` column is 21,183 km. That's more than half the circumference of the earth!

```
In [27]: # Calculate median values of all columns for churned and retained users
churned_df = df[df['label'] == 'churned']
retained_df = df[df['label'] == 'retained']

churned_median = churned_df.median()
retained_median = retained_df.median()

# Print the median values
print("Median values for churned users:")
print(churned_median)

print("\nMedian values for retained users:")
print(retained_median)
```

Median values for churned users:

ID	7477.500000
sessions	59.000000
drives	50.000000
total_sessions	164.339042
n_days_after_onboarding	1321.000000
total_navigations_fav1	84.500000
total_navigations_fav2	11.000000
driven_km_drives	3652.655666
duration_minutes_drives	1607.183785
activity_days	8.000000
driving_days	6.000000
dtype:	float64

Median values for retained users:

ID	7509.000000
sessions	56.000000
drives	47.000000
total_sessions	157.586756
n_days_after_onboarding	1843.000000
total_navigations_fav1	68.000000
total_navigations_fav2	9.000000
driven_km_drives	3464.684614
duration_minutes_drives	1458.046141
activity_days	17.000000
driving_days	14.000000
dtype:	float64

This offers an interesting snapshot of the two groups, churned vs. retained:

Users who churned averaged ~3 more drives in the last month than retained users, but retained users used the app on over twice as many days as churned users in the same time period.

The median churned user drove ~200 more kilometers and 2.5 more hours during the last month than the median retained user.

It seems that churned users had more drives in fewer days, and their trips were farther and longer in duration. Perhaps this is suggestive of a user profile. Continue exploring!

Calculate the median kilometers per drive in the last month for both retained and churned users.

Begin by dividing the `driven_km_drives` column by the `drives` column. Then, group the results by churned/retained and calculate the median km/drive of each group.

```
In [28]: # Add a column to df called `km_per_drive`

df['km_per_drive'] = df['driven_km_drives'] / df['drives']

# Group by `label`, calculate the median, and isolate for km per drive
median_km_per_drive_by_label = df.groupby('label')['km_per_drive'].median()

print(median_km_per_drive_by_label)
```

```
label
churned    74.109416
retained   75.014702
Name: km_per_drive, dtype: float64
```

The median retained user drove about one more kilometer per drive than the median churned user. How many kilometers per driving day was this?

To calculate this statistic, repeat the steps above using `driving_days` instead of `drives`.

```
In [29]: # Add a column to df called `km_per_driving_day`
df['km_per_driving_day'] = df['driving_days'] / df['drives']

# Group by `label`, calculate the median, and isolate for km per driving day
km_per_driving_day = df.groupby('label')['km_per_driving_day'].median()

print(km_per_driving_day)
```

```
label
churned    0.100000
retained    0.246154
Name: km_per_driving_day, dtype: float64
```

Now, calculate the median number of drives per driving day for each group.

```
In [31]: # Add a column to df called `drives_per_driving_day`
df['drives_per_driving_day'] = df['drives'] / df['driving_days']

# Group by `label`, calculate the median, and isolate for drives per driving
drives_per_driving_day = df.groupby('label')['drives_per_driving_day'].median()
print(drives_per_driving_day)
```

```
label
churned      10.0000
retained      4.0625
Name: drives_per_driving_day, dtype: float64
```

The median user who churned drove 698 kilometers each day they drove last month, which is almost ~240% the per-drive-day distance of retained users. The median churned user had a similarly disproportionate number of drives per drive day compared to retained users.

It is clear from these figures that, regardless of whether a user churned or not, the users represented in this data are serious drivers! It would probably be safe to assume that this data does not represent typical drivers at large. Perhaps the data—and in particular the sample of churned users—contains a high proportion of long-haul truckers.

In consideration of how much these users drive, it would be worthwhile to recommend to Waze that they gather more data on these super-drivers. It's possible that the reason for their driving so much is also the reason why the Waze app does not meet their specific set of needs, which may differ from the needs of a more typical driver, such as a commuter.

Finally, examine whether there is an imbalance in how many users churned by device type.

Begin by getting the overall counts of each device type for each group, churned and retained.

```
In [32]: # For each label, calculate the number of Android users and iPhone users
# For each label, calculate the number of Android users and iPhone users
android_users_by_label = df[df['device'] == 'Android'].groupby('label').size()
iphone_users_by_label = df[df['device'] == 'iPhone'].groupby('label').size()
```

Now, within each group, churned and retained, calculate what percent was Android and what percent was iPhone.

```
In [33]: # For each label, calculate the percentage of Android users and iPhone users
total_users_by_label = df.groupby('label').size()
android_percentage_by_label = (android_users_by_label / total_users_by_label)
iphone_percentage_by_label = (iphone_users_by_label / total_users_by_label)
```

```

print("\nAndroid users by label:")
print(android_users_by_label)

print("\niPhone users by label:")
print(iphone_users_by_label)

print("\nAndroid percentage by label:")
print(android_percentage_by_label)

print("\niPhone percentage by label:")
print(iphone_percentage_by_label)

```

Android users by label:

```

label
churned      891
retained    4183
dtype: int64

```

iPhone users by label:

```

label
churned      1645
retained    7580
dtype: int64

```

Android percentage by label:

```

label
churned      35.134069
retained     35.560656
dtype: float64

```

iPhone percentage by label:

```

label
churned      64.865931
retained     64.439344
dtype: float64

```

The ratio of iPhone users and Android users is consistent between the churned group and the retained group, and those ratios are both consistent with the ratio found in the overall dataset.